

## ▼ Bank\_Customer\_churn\_Prediction using Deep Learning

```
#Importing Required Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error as mse, mean_absolute_error as mae, confusion_matrix as cm
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

import keras
import warnings
warnings.filterwarnings('ignore')

from google.colab import drive
drive.mount('/content/drive')
```



Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive').

```
#import data
df=pd.read_csv('/content/drive/MyDrive/DL /Churn_Modelling.csv')
```

### EDA - Exploratory Data Analysis

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1
2	3	15619304	Onio	502	France	Female	42	8
3	4	15701354	Boni	699	France	Female	39	1
4	5	15737888	Mitchell	850	Spain	Female	43	2

```
df.tail()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
9995	9996	15606229	Obijaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	10
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	

df.shape

(10000, 14)

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender                10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                10000 non-null  int64
8   Balance               10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

df.describe()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Ba
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.0
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.8
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.4
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.0
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.0
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.5
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.2
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.0

```
# Check columns list and missing values
df.isnull().sum()
```

```

CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
Exited         int64
dtype: object


```

```

plt.figure(figsize=(12,10))
sns.heatmap(df.corr(), annot = True)

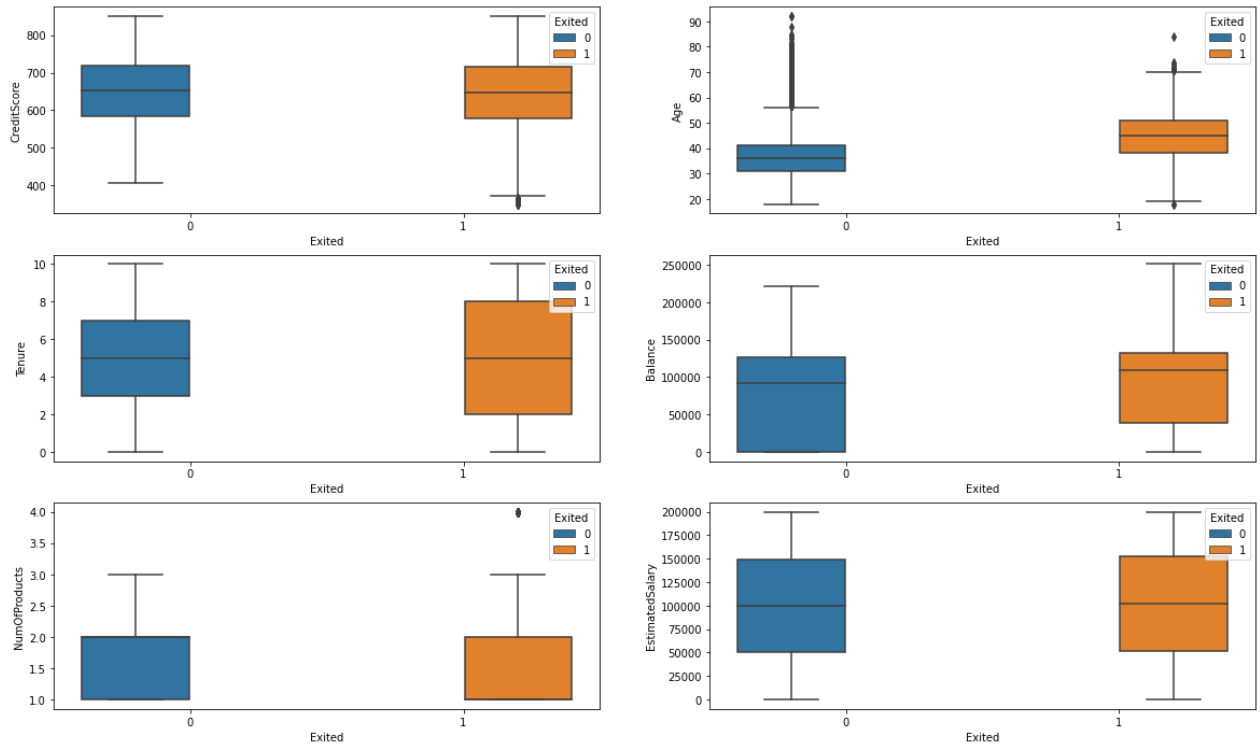
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f535b14a3d0>



```
fig, axarr = plt.subplots(3, 2, figsize=(20, 12))
sns.boxplot(y='CreditScore',x = 'Exited', hue = 'Exited',data = df, ax=axarr[0][0])
sns.boxplot(y='Age',x = 'Exited', hue = 'Exited',data = df , ax=axarr[0][1])
sns.boxplot(y='Tenure',x = 'Exited', hue = 'Exited',data = df, ax=axarr[1][0])
sns.boxplot(y='Balance',x = 'Exited', hue = 'Exited',data = df, ax=axarr[1][1])
sns.boxplot(y='NumOfProducts',x = 'Exited', hue = 'Exited',data = df, ax=axarr[2][0])
sns.boxplot(y='EstimatedSalary',x = 'Exited', hue = 'Exited',data = df, ax=axarr[2][1])
```







<matplotlib.axes.\_subplots.AxesSubplot at 0x7f535aae1750>




Check numerical data distribution

```
num_cols = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
df.hist(column=num_cols, figsize=(14,10))
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f5359fb21d0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f535b234650>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f535abcd450>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f535aba5950>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f5361860e50>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f5361824390>]],
      dtype=object)
```



```
# Handling Outliers
fig, ax = plt.subplots(6, 1, figsize=(18,24))
for i in range(6):
    sns.boxplot(x = df[num_cols[i]], ax=ax[i])
```



## Feature Engineering

```
#Drop identifier data column
df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis = 1)

#Move outliers values to the upper and lower bounds
for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    S = 1.5*IQR
    LB = Q1 - S
    UB = Q3 + S
    df.loc[df[col] > UB, col] = UB
    df.loc[df[col] < LB, col] = LB

#Create new Gender column with value is 0 and 1
df['Gender New'] = pd.factorize(df.Gender)[0]
df = df.drop(['Gender'], axis = 1)

#One hot encode Geography column
dvcats_dummies = pd.get_dummies(df.Geography)
df=pd.concat([df, dvcats_dummies], axis=1)

df = df.drop(['Geography'], axis = 1)

df.head(10)
```



CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
-------------	-----	--------	---------	---------------	-----------	----------------

```
# split feature, targer and train, test
```

```
X = df.drop(columns=['Exited'])
```

```
y = df['Exited'].values
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=0)
```

4	850	43	2	12551082	1 0	1	1
---	-----	----	---	----------	-----	---	---

```
#Feature scaling
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_trf = scaler.fit_transform(X_train)
```

```
X_test_trf = scaler.transform(X_test)
```

```
le = LabelEncoder()
```

```
y_train = le.fit_transform(y_train)
```

```
y_test = le.fit_transform(y_test)
```

```
y_train
```

```
array([0, 0, 0, ..., 0, 0, 1])
```

```
y_test
```

```
array([0, 1, 0, ..., 0, 0, 0])
```

## Model Building

```
# Neural network
```

```
model = Sequential()
```

```
model.add(Dense(6, activation = 'relu', input_shape=(12,)))
```

```
model.add(keras.layers.Dropout(0.5))
```

```
model.add(Dense(6, activation = 'relu'))
```

```
model.add(keras.layers.Dropout(0.5))
```

```
model.add(Dense(1, activation = 'sigmoid'))
```

```
model.compile(optimizer = 'Adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

## Model training

```
history = model.fit(X_train,y_train,batch_size=50,epochs=100,verbose=1,validation_split=0.
```


```
Epoch 1/100
```



```
128/128 [=====] - 1s 5ms/step - loss: 11241.8672 - accuracy: 0.0000  
Epoch 2/100  
128/128 [=====] - 0s 3ms/step - loss: 3341.2812 - accuracy: 0.0000  
Epoch 3/100  
128/128 [=====] - 0s 3ms/step - loss: 1255.6195 - accuracy: 0.0000  
Epoch 4/100  
128/128 [=====] - 0s 3ms/step - loss: 561.2086 - accuracy: 0.0000  
Epoch 5/100  
128/128 [=====] - 0s 3ms/step - loss: 237.1535 - accuracy: 0.0000  
Epoch 6/100  
128/128 [=====] - 0s 3ms/step - loss: 81.8563 - accuracy: 0.0000  
Epoch 7/100  
128/128 [=====] - 1s 5ms/step - loss: 10.8184 - accuracy: 0.0000  
Epoch 8/100  
128/128 [=====] - 1s 6ms/step - loss: 8.1663 - accuracy: 0.0000  
Epoch 9/100  
128/128 [=====] - 1s 5ms/step - loss: 4.1229 - accuracy: 0.0000  
Epoch 10/100  
128/128 [=====] - 1s 6ms/step - loss: 1.9319 - accuracy: 0.0000  
Epoch 11/100  
128/128 [=====] - 1s 5ms/step - loss: 0.9040 - accuracy: 0.0000  
Epoch 12/100  
128/128 [=====] - 1s 5ms/step - loss: 0.6583 - accuracy: 0.0000  
Epoch 13/100  
128/128 [=====] - 1s 5ms/step - loss: 0.6699 - accuracy: 0.0000  
Epoch 14/100  
128/128 [=====] - 1s 7ms/step - loss: 0.5520 - accuracy: 0.0000  
Epoch 15/100  
128/128 [=====] - 1s 9ms/step - loss: 0.5662 - accuracy: 0.0000  
Epoch 16/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5433 - accuracy: 0.0000  
Epoch 17/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5489 - accuracy: 0.0000  
Epoch 18/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5841 - accuracy: 0.0000  
Epoch 19/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5232 - accuracy: 0.0000  
Epoch 20/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5101 - accuracy: 0.0000  
Epoch 21/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5131 - accuracy: 0.0000  
Epoch 22/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5149 - accuracy: 0.0000  
Epoch 23/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5084 - accuracy: 0.0000  
Epoch 24/100  
128/128 [=====] - 0s 4ms/step - loss: 0.5058 - accuracy: 0.0000  
Epoch 25/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5123 - accuracy: 0.0000  
Epoch 26/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5181 - accuracy: 0.0000  
Epoch 27/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5079 - accuracy: 0.0000  
Epoch 28/100  
128/128 [=====] - 0s 3ms/step - loss: 0.5054 - accuracy: 0.0000  
Epoch 29/100
```

plt.figure()

```
plt.plot(history.history['loss'])  
plt.plot(history.history['accuracy'])  
plt.show()
```




```
#Model evaluation
```

```
y_pred = model.predict(X_test)  
y_pred = (y_pred > 0.5)
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])
```

```
[<matplotlib.lines.Line2D at 0x7f5365332190>]
```




```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])
```

[<matplotlib.lines.Line2D at 0x7f53652e7250>]



```
#Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8,8))
sns.heatmap(cm, annot=True, fmt=".0f",linewidth=.5, square=True);
plt.xlabel('Actual Label');
plt.ylabel('Predicted Label');
```



```
#Classification report
score = metrics.accuracy_score(y_test,y_pred)
print("Accuracy:", score)

print("Report:",metrics.classification_report(y_test,y_pred))
```

Accuracy: 0.7975

Report:                      precision      recall      f1-score      support

0	0.80	1.00	0.89	1595
1	0.00	0.00	0.00	405

accuracy			0.80	2000
macro avg	0.40	0.50	0.44	2000

weighted avg	0.64	0.80	0.71	2000
--------------	------	------	------	------


## MAE, MSE & RMSE

```
MSE = mse(y_test, y_pred)
MAE = mae(y_test, y_pred)
print("MSE:",MSE)
print("RMSE:",np.sqrt(MSE))
print("MAE:",MAE)
```

```
MSE: 0.2025
RMSE: 0.45
MAE: 0.2025
```


## Model loss

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



## Model accuracy

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



[Colab paid products](#) - [Cancel contracts here](#)

✓ 3s completed at 11:56

